




Common SQL Anti-Patterns and How to Avoid Them

The
Pragmatic
Programmers

SQL Antipatterns

Avoiding the Pitfalls of Database Programming



Bill Karwin
Edited by Jacquelyn Carter

 Search ...

RECENT POSTS

[How to Implement a GraphQL Federation API for Microservices](#)

[How to Build an API for AI-Powered Text Summarization](#)

[How to Use the IBM Watson API for AI and Machine Learning Applications](#)

[How to Use the Payoneer API for Payment Processing](#)

[How to Implement API Key Expiration Policies for Better Security](#)

Sure! Common SQL anti-patterns are common mistakes or bad practices that developers may fall into when writing SQL queries. These anti-patterns can lead to performance issues, inefficiencies, and security vulnerabilities in the database. In order to avoid these pitfalls, developers should be aware of these anti-patterns and follow best practices when writing SQL queries. In this guide, we will explore some of the most common SQL anti-patterns and provide tips on how to avoid them for better database performance and security.

In SQL programming, it's crucial to write clean, efficient, and maintainable code. However, many developers fall into common SQL **anti-patterns** that can lead to performance issues, data integrity problems, and scalability challenges. Understanding these anti-patterns and learning how to avoid them is essential for anyone who interacts with relational databases. In this article, we will explore various SQL anti-patterns along with strategies to circumvent them.

Table of Contents



- 1. 1. SELECT * Anti-Pattern
- 2. 2. Lack of Indexing
- 3. 3. Inappropriate Use of Joins
- 4. 4. Not Normalizing Data
- 5. 5. Using Cursors When Not Needed
- 6. 6. Neglecting Transactions
- 7. 7. Improper Use of NULLs
- 8. 8. Poor Naming Conventions
- 9. 9. Hardcoding Values
- 10. 10. Overusing Subqueries
- 11. 11. Ignoring Error Handling
- 12. 12. Ignoring Best Practices in SQL Development
- 12.1. Related posts:

1. SELECT * Anti-Pattern

Using **SELECT *** in your SQL queries seems convenient but can lead to several issues:

How to Use the TikTok API for Content Automation

How to Use the Expedia API for Travel Booking Services

How to Implement API Dependency Resolution in Microservices

How to Build an API That Supports File Streaming

How to Use the Amazon Alexa API for Voice Command Integrations

- **Performance Degradation:** Fetching all columns can slow down your queries, especially when dealing with large datasets.
- **Excessive Data Transfer:** Returning more data than necessary increases load on the network and the client application.
- **Schema Changes:** If the database schema changes, your application might break unexpectedly.

How to Avoid: Always specify the exact columns you need in your SELECT statements:

```
SELECT column1, column2 FROM your_table WHERE conditi
```



2. Lack of Indexing

Failing to create **indexes** on frequently queried columns can lead to poor database performance:

- **Slow Queries:** Without indexes, the database must perform full table scans, which are resource-intensive and time-consuming.
- **Negative User Experience:** Long-running queries can lead to timeouts and frustration for end-users.

How to Avoid: Analyze query performance and implement indexes strategically:

```
CREATE INDEX idx_column_name ON your_table (column_na
```



3. Inappropriate Use of Joins

Joins can significantly enhance data retrieval but misusing them can cause serious performance issues:

- **Cartesian Products:** Forgetting to include join conditions can result in massive datasets being returned, severely impacting performance.
- **Multiple Joins:** Too many joins in a single query can make it complex and slow.

How to Avoid: Always use proper join conditions and consider breaking down complex queries into simpler parts:

```
SELECT a.column1, b.column2 FROM table_a a  
JOIN table_b b ON a.id = b.a_id;
```

4. Not Normalizing Data

Data normalization is crucial for maintaining **data integrity** and reducing redundancy. Failure to normalize can lead to:

- **Update Anomalies:** Duplicated data can cause inconsistencies during updates.
- **Insertion Anomalies:** Difficulty adding new data without redundant information.

How to Avoid: Follow the rules of normalization and aim for at least third normal form (3NF) for your database schema:

Break down larger tables into smaller, related tables and establish **foreign keys** to maintain relationships.

5. Using Cursors When Not Needed

Cursors can be convenient for row-by-row processing, but they are generally less efficient than set-based operations:

- **Performance Issues:** Cursors can create significant overhead and slow down operations.
- **Complexity:** Cursors can make your code more complicated and harder to maintain.

How to Avoid: Use set-based operations instead of cursors whenever possible:

```
UPDATE your_table  
SET column_name = new_value  
WHERE condition;
```

6. Neglecting Transactions

Not using transactions when performing multiple related operations can jeopardize data consistency:

- **Partial Updates:** If one operation fails, previous operations may leave the database in an inconsistent state.
- **Data Loss:** Failure to rollback changes can result in lost data due to errors.

How to Avoid: Always use transactions when multiple changes depend on each other:

```
BEGIN TRANSACTION;  
-- multiple SQL statements  
COMMIT; -- or ROLLBACK in case of error
```

7. Improper Use of NULLs

Using **NULL** values can lead to ambiguity and unexpected behavior in SQL:

- **Complicated Logic:** Queries can become difficult to interpret and debug.
- **Performance Hits:** NULL values can complicate indexing and slow down performance.

How to Avoid: Use **NOT NULL** constraints where possible and consider default values:

```
CREATE TABLE your_table (
    id INT NOT NULL,
    name VARCHAR(255) NOT NULL DEFAULT 'Unknown'
);
```

8. Poor Naming Conventions

Using vague and inconsistent naming conventions for tables and columns can lead to confusion:

- **Maintenance Challenges:** Understanding the schema becomes difficult.
- **Collaboration Issues:** Other developers may struggle to work with poorly named structures.

How to Avoid: Establish clear and consistent naming conventions that convey meaning. Use prefixes or suffixes that describe the content:

```
CREATE TABLE customer_data (
    customer_id INT,
```

```
    first_name VARCHAR(50),  
    last_name VARCHAR(50)  
);
```

9. Hardcoding Values

Hardcoding values directly in SQL queries can make your code inflexible and error-prone:

- **Security Risks:** Hardcoded data can expose sensitive information and increase vulnerability.
- **Maintenance Burden:** Making changes requires altering the codebase rather than just the data.

How to Avoid: Use parameters in your queries to enhance flexibility and security:

```
SELECT * FROM your_table WHERE column_name = ?;
```

10. Overusing Subqueries

Subqueries can greatly simplify SQL logic. However, excessive use can lead to performance issues:

- **Execution Time:** Nested subqueries are often slower than joined queries.
- **Readability Issues:** Complex queries can become hard to read and maintain.

How to Avoid: Analyze the necessity of subqueries versus joins and opt for joins when appropriate:

```
SELECT a.column1, b.column2 FROM table_a a  
JOIN (SELECT column2 FROM table_b WHERE condition) b 0
```



11. Ignoring Error Handling

Not properly addressing potential errors can cause applications to fail silently, leading to frustrating user experiences:

- **Undetected Problems:** Lack of error handling can result in data corruption or loss.
- **Difficulty Troubleshooting:** Issues become hard to track without proper logging or escalation of errors.

How to Avoid: Implement robust error handling and logging mechanisms to capture and respond to issues:

```
BEGIN TRY  
    -- Your SQL code here  
END TRY  
  
BEGIN CATCH  
    -- Log error details  
END CATCH;
```

12. Ignoring Best Practices in SQL Development

Overlooking SQL development best practices can result in inefficient and unproductive database management:

- **Lack of Documentation:** Failing to document your SQL code makes it difficult for future developers to understand it.
- **Poor Query Optimization:** Not reviewing queries for optimization can lead to subpar performance.

How to Avoid: Adhere to SQL coding standards, regularly review your code, and document every change:

Finally, stay updated with the latest trends and techniques in SQL and relational databases to continually improve your database skills and maintain high standards of database management.

Common SQL Anti-Patterns can cause inefficiencies and errors in database queries. By understanding these pitfalls and implementing best practices such as proper indexing, normalization, and avoiding unnecessary database calls, developers can significantly improve the performance and reliability of their SQL queries. By being mindful of these anti-patterns and adopting good practices, developers can avoid common mistakes and optimize their database interactions effectively.

Related Posts:



[The SELECT Statement: How to](#)

```
mysql> SELECT * FROM student_marks
ORDER BY CASE
WHEN stu_branch = 'C.S.E' then 1
WHEN stu_branch = 'I.T' then 2
WHEN stu_branch = 'E.C.E' then 3
WHEN stu_branch = 'I.C.E' then 4
WHEN stu_branch = 'E.E.E' then 5
END ASC;

```

Results (10 rows)

No.	Stu_name	Stu_branch	Max_Marks
2001	Rakesh	C.S.E	570
2002	Sumeet	C.S.E	600
2003	Prashant	C.S.E	580
2004	Shivam	I.T	560
2005	Naman	I.T	510
2006	Pradeep	E.C.E	590
2007	Yash	E.C.E	540
2008	Pratik	E.C.E	560
2009	Harsh	I.C.E	580


[How to Use ORDER BY to Sort Query](#)



[The DELETE Command:](#)




[Understanding Primary Keys and](#)



Common Data Types in SQL Explained




The CASE Statement in SQL: Conditional




Limiting Results with the LIMIT

Operator Symbol	Functionality
and	If both the operands are true, then the condition becomes true.
or	If any of the two operands are true, then the condition becomes true.
not	Used to reverse the logical state of its operand.


Logical Operators: AND, OR, NOT



Date Functions: NOW(), DATEDIFF(),



What are Window Functions, and How



Understanding ROW_NUMBER(), RANK()


Pivot

Rep	Sales	row_number()	rank()
Haris	10	1	1
Alex	7	2	2
John	7	3	3
Bryan	6	4	4
James	2	5	5
Karen	1	6	6

Unpivot

Country	Year
USA	2020
USA	2021
France	2020
France	2021
Germany	2020
Germany	2021

PIVOT and UNPIVOT Explained with



Types of Indexes: Clustered vs. Non-

IDE ANALYZE processing...
note an index on "convicts"(“last_name” text,pattern_upper) doesn't include 'image' in the SELECT list or add LIMIT. I used ORDER BY in the subquery. Note that DISTINCT. The primary key always is.
I wonder if you have ORDER BY in a subquery?
is the index “convicts_releasemt_idx”. It hasn't been in ages.
Just “random_page_cost” to 1.2 to match your actual
indexing needs. That is pretty fast.
INDEX INDEX “convicts_sentence_idx”, it is already
atty bloated.

IDE -

Using EXPLAIN to Analyze SQL Queries

PREVIOUS

Choosing Between AWS, Azure,
and Google Cloud for SQL
Databases

NEXT

Data Compression for
Relational Databases

Leave a Reply

Your email address will not be published. Required fields are marked *

COMMENT *

NAME *

EMAIL *

WEBSITE

Save my name, email, and website in this browser for the next time I comment.

POST COMMENT

Copyright © 2025 Datatas.